

STOR566: Introduction to Deep Learning

Lecture 12: Generative Models I

Yao Li
UNC Chapel Hill


Oct 10, 2024

Unsupervised Learning

- Working with datasets without a **response** variable
- Some Applications:
 - Clustering
 - Data Compression
 - Exploratory Data Analysis
 - Generating New Examples
 - ...
- Example: PCA, K-means, Autoencoders, GAN, etc


Autoencoder: Basic Architecture

- Autoencoder: A special type of DNN where the target (response) of each input is the input itself.



Autoencoder: Basic Architecture

- Autoencoder: A special type of DNN where the target (response) of each input is the input itself.



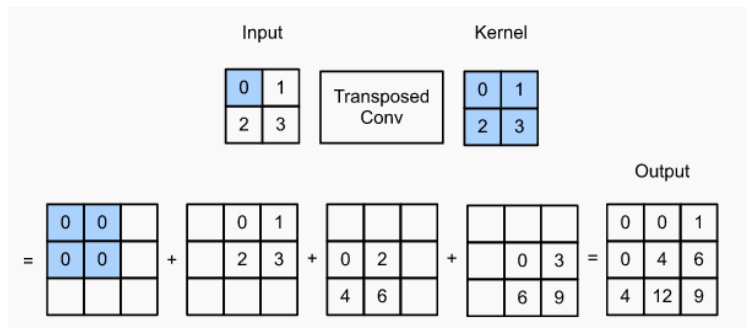
- Objective:

$$\|x - D(E(x))\|^2$$

Encoder: $E : \mathbb{R}^n \rightarrow \mathbb{R}^d$

Decoder: $D : \mathbb{R}^d \rightarrow \mathbb{R}^n$


Transposed Convolution



(Figure from Dive into Deep Learning)

- Multiple input and output channels: works the same as the regular convolution
- Number of weights: $k_1 \times k_2 \times d_{in} \times d_{out} + d_{out}$

Transposed Convolution



(Figure from Dive into Deep Learning)

- Strides are specified for the output feature map
- Padding: remove rows and columns from the output

Overfitting

- Overfitting is a problem
- Solutions:
 - Bottleneck layer: a low-dimensional representation of the data ($d < n$)
 - Denoise autoencoder
 - Sparse autoencoder
 - ...

Regularization

- Objective:

$$L(\mathbf{x}, \hat{\mathbf{x}}) + \text{regularizer},$$

Regularization

- Objective:

$$L(\mathbf{x}, \hat{\mathbf{x}}) + \text{regularizer},$$

$L(\cdot, \cdot)$: captures the distance between the input (\mathbf{x}) and the output ($\hat{\mathbf{x}}$).

- Example: $\|\mathbf{x} - \hat{\mathbf{x}}\|^2$

Regularization

- Objective:

$$L(\mathbf{x}, \hat{\mathbf{x}}) + \text{regularizer},$$


$L(\cdot, \cdot)$: captures the distance between the input (\mathbf{x}) and the output ($\hat{\mathbf{x}}$).

- Example: $\|\mathbf{x} - \hat{\mathbf{x}}\|^2$

Regularizer example:

- L_1 penalty: $\sum_j |h_j^l|$
- h_j^l : hidden output of j -th neuron in l -th layer

Sparse Autoencoder

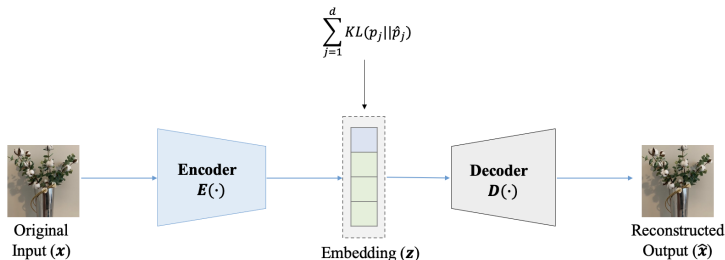


- Objective:

$$\|\mathbf{x} - \mathbf{D}(\mathbf{E}(\mathbf{x}))\|^2 + \lambda \sum_j |z_j|$$

- Iterate over layers.

Sparse Autoencoder



- Another regularizer:

$$\|x - D(E(x))\|^2 + \lambda \sum_j KL(p_j || \hat{p}_j)$$

- Convert value of z to $[0, 1]$. (e.g., sigmoid activation)
- p_j : probability of activation for neuron j in the bottleneck layer
- $\hat{p}_j = \frac{1}{B} \sum_{i=1}^B z_{ij}$

Denoising Autoencoder





Figure from Bank, Dor, Noam Koenigstein, and Raja Giryes. "Autoencoders." (2020).

- Another regularizer:

$$\|x - D(E(x + \delta))\|^2$$

- δ : Random noise

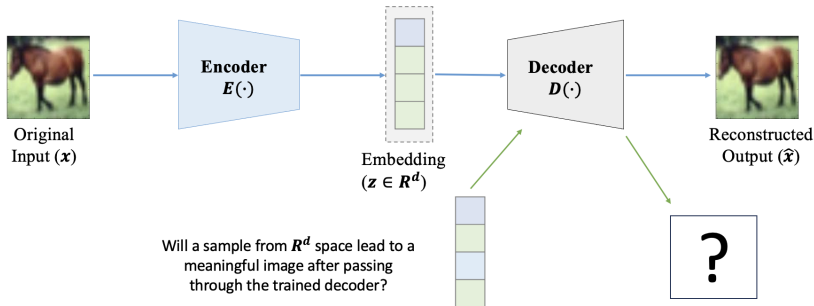
Denoising Autoencoder



- noisy data → clean data
- Learn to capture valuable features and ignore noise


Generative Model

Generative Problem





- In general, a trained Vanilla auto-encoder cannot be used to generate new data

Variational Autoencoder (VAE)



- Probabilistic model: will let us generate data from the model
- Encoder outputs μ and σ
- Draw $\tilde{z} \sim N(\mu, \sigma)$
- Decoder decodes this **latent** variable \tilde{z} to get the output

Variational Autoencoder (VAE)



- Maximum likelihood approach: $\prod_i p(\mathbf{x}_i)$
- Variational lower bound as objective:
 - End-to-End reconstruction loss (e.g., square loss)
 - Regularizer: $KL(q_\phi(z|x)||p(z))$
- Objective:

$$L(\mathbf{x}, \hat{\mathbf{x}}) + KL(q_\phi(z|x)||p(z))$$

Re-parameterization Trick




Figure from Jeremy Jordon Blog

- Cannot back-propagate error through random samples
- Reparameterization trick: replace $\tilde{z} \sim N(\mu, \sigma)$ with $\epsilon \sim N(0, I)$,
 $z = \epsilon\sigma + \mu$

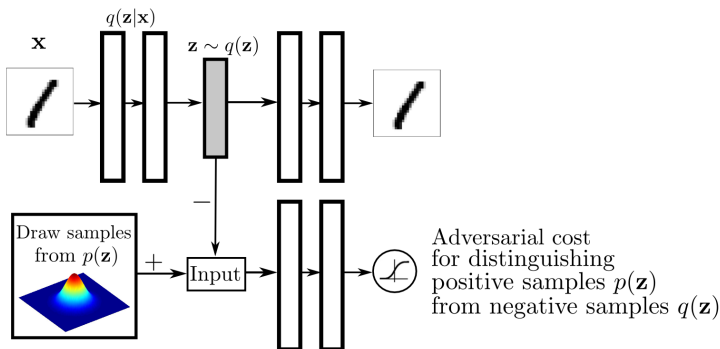
Variational Lower Bound

- Variational lower bound:

$$\log p(x) \geq E_{q(z|x)} (\log p(x|z)) + KL(q(z|x)||p(z))$$

- How to derive the variational lower bound from the likelihood?

Adversarial Autoencoder





- The top row is a standard autoencoder
- Force the embedding space distribution towards the prior

Embedding Space Visualization

Commonly used visualization tools:

- t-SNE (t-Distributed Stochastic Neighbor Embedding)
 - Van der Maaten et al. (2008). Visualizing data using t-SNE. Journal of Machine Learning Research, 9(11).
 - Available: sklearn
- UMAP (Uniform Manifold Approximation and Projection)
 - McInnes et al. (2018). UMAP: Uniform Manifold Approximation and Projection. Journal of Open Source Software, 3(29), 861,
 - Available: umap-learn
- PCA (Principal Component Analysis)
 - Available: sklearn

Examples with tSNE



- Embedding space visualization for a Vanilla autoencoder and a VAE trained on MNIST
- VAE: more compact

Examples with PCA

- Problem: Game Result Prediction



Figure: Heroes of the Storm and Dota 2 characters

Assumption

Assumption

We assume a team's score can be written as

$$s_t^+ = \sum_{i \in I_t^+} w_i + \sum_{i \in I_t^+} \sum_{j \in I_t^+} \mathbf{v}_i^T \mathbf{v}_j$$

- w_i : individual ability of i -th player
- $\mathbf{v}_i \in R^d$: teamwork ability of i -th player
- I_t^+ : winning team player index set
- s_t^+ : winning team score

Team Ability Visualization (PCA)

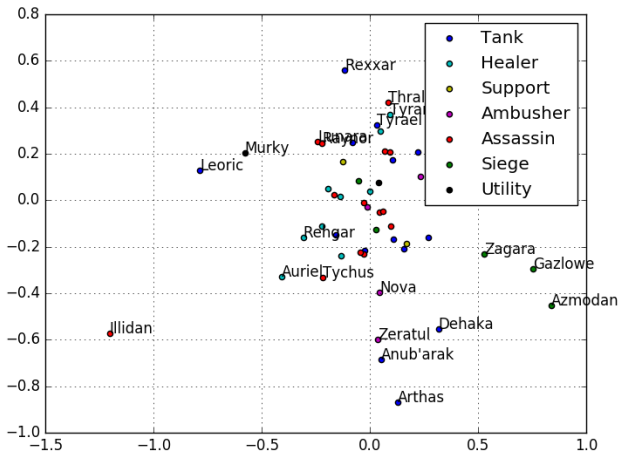


Figure: Projection of team ability vector for each character (\mathbf{v}_i) to 2-D space. Colors represents the official categorization for these characters.

Conclusions

- Autoencoder
- Regularization
- Variational Autoencoder
- Visualization tools

Questions?